



CS 590 – Algorithms

Lecture 7 – Red-Black Trees



Outlines

7. Red-black trees (RBTs)

7.1. Characteristics and Properties

7.2. Black Heights

7.3. Operations

7.3.1. Rotations

7.3.2. Insertion

7.3.3. Deletion



7. Red-black trees

Red-black trees

- A variation of binary search trees.
- Self-balancing BST:
 - *Balanced*: height is $O(\lg n)$, where n is the number of nodes.
 - Operations will take $O(\lg n)$ time in the worst case.

7.1. Properties



Red-black trees:

- All attributes of BST will be inherited in ***red-black tree (RBT)***
- But RBT will have one extra (+1) bit per node.
 - New attribution: `x.color` - either **red** or **black**.
- RBT uses a single sentinel, `T.nil`, for all the leaves.
- `T.nil.color = black`
- `T.root.p = T.nil`



7.1. Properties

RBT Color Properties:

1. Every node, $x.\text{color}$, is either red or black.
2. $T.\text{root}.\text{color} = \text{black}$.
3. $T.\text{nil}.\text{color} = \text{black}$.
4. If $x.\text{color} = \text{red}$, $x.\text{left}.\text{color} \neq \text{red}$ and $x.\text{right}.\text{color} \neq \text{red}$.
5. All paths from x to its descent will have the **same number of black nodes** (see black heights, $\text{bh}(x)$).

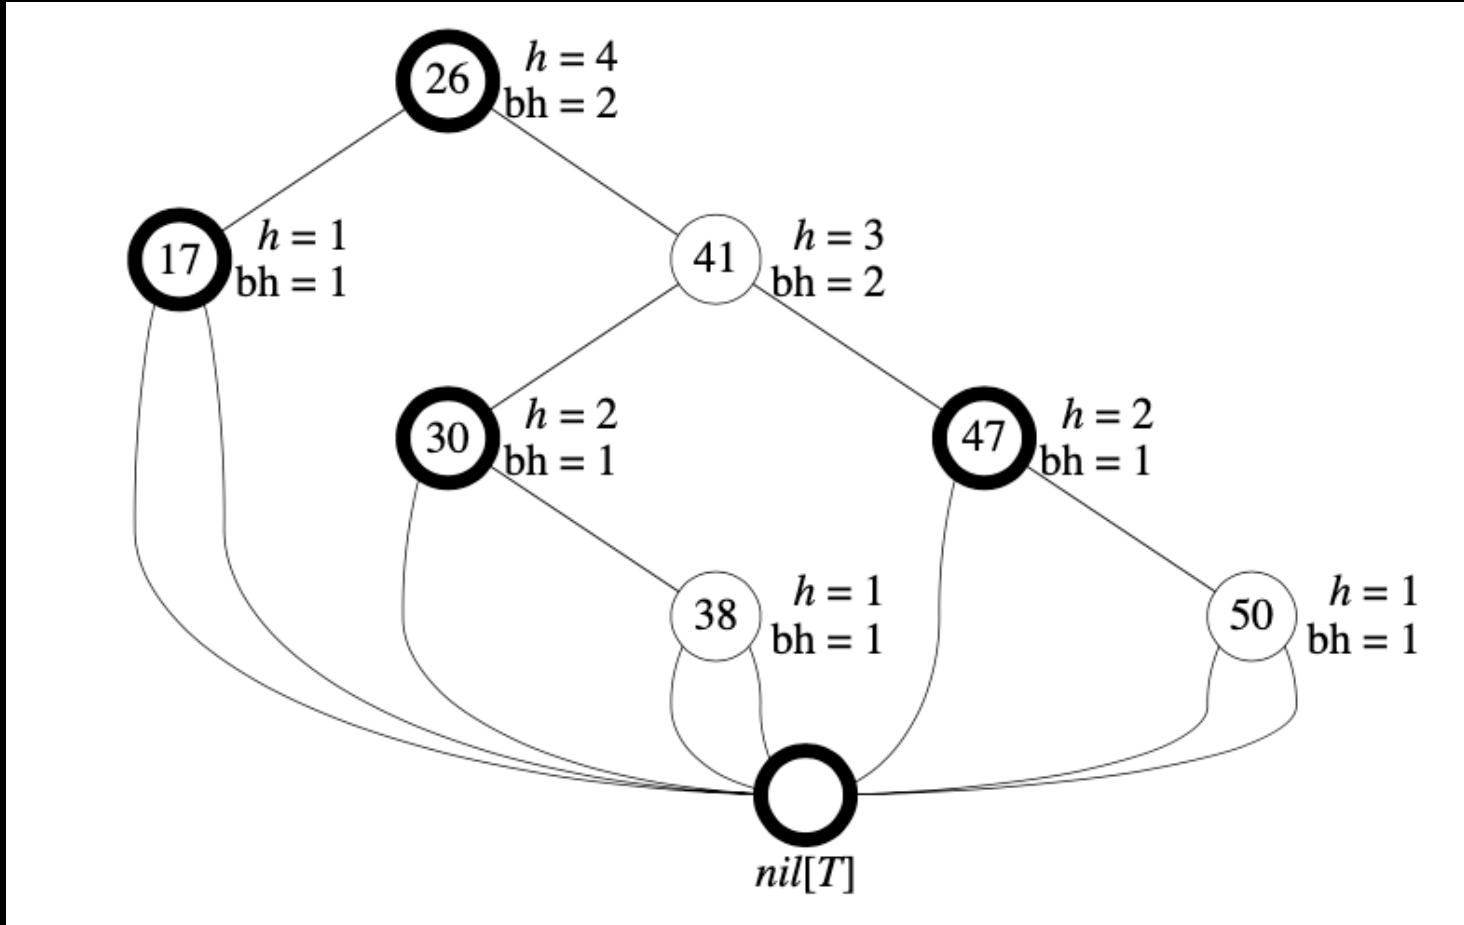


7.1. Properties

Black Height (bh) Properties:

- The height of the tree, $h(T)$, is the number of edges in the longest path to a leaf.
- The black height, $bh(x)$, is the number of black nodes from a node x to a leaf.
 - $x.\text{color}$ is exclusive.
 - $T.\text{nil}$ is inclusive.

7.2. Black Heights





7.2. Black Heights

- **Claim 1:**
 - If $h(x) = h$, $bh(x)$ is $\geq h/2$ for all x in T .
- **Claim 2:**
 - A subtree of x contains $\geq 2^{bh(x)} - 1$ internal nodes.
- **Lemma 1:**
 - An RBT with n -many internal nodes has $h \leq 2 * \lg(n+1)$.



7.3. Operations

Operations on red-black trees

- The non-modifying BST operations will be carried over.
 - Inorder-Tree-Walk, Minimum, Successor...
 - Except Insertion and Deletion.
- All searching algorithms will have the running time in $O(h)$.



7.3.1. Rotations

Rotation Operation:

- is the basic tree-restructuring operation taking a node x within T .
- is to maintain RBTs as self-balanced BSTs.
- will use pointers to change the local pointer structure.
- must not upset the BST's property.
- operates in two directions, left and right.
- separate implementation is easier.



7.3.1. Rotations

LEFT-ROTATE(T, x) algorithm implementation.

1. Set y as $x.\text{right}$:
2. Turn the left subtree of y into the right of x :
3. What happens if $y.\text{left}$ is not nil?
 - Link point x as $y.\text{left}.\text{p}$
4. Link $x.\text{p}$ to $y.\text{p}$
5. What if $x.\text{p} = T.\text{nil}$?
 - What becomes the root of the tree?
 - What if x is the left child?
 - What if x is the right child?



7.3.1. Rotations

Algorithm (LEFT-ROTATE(T, x))

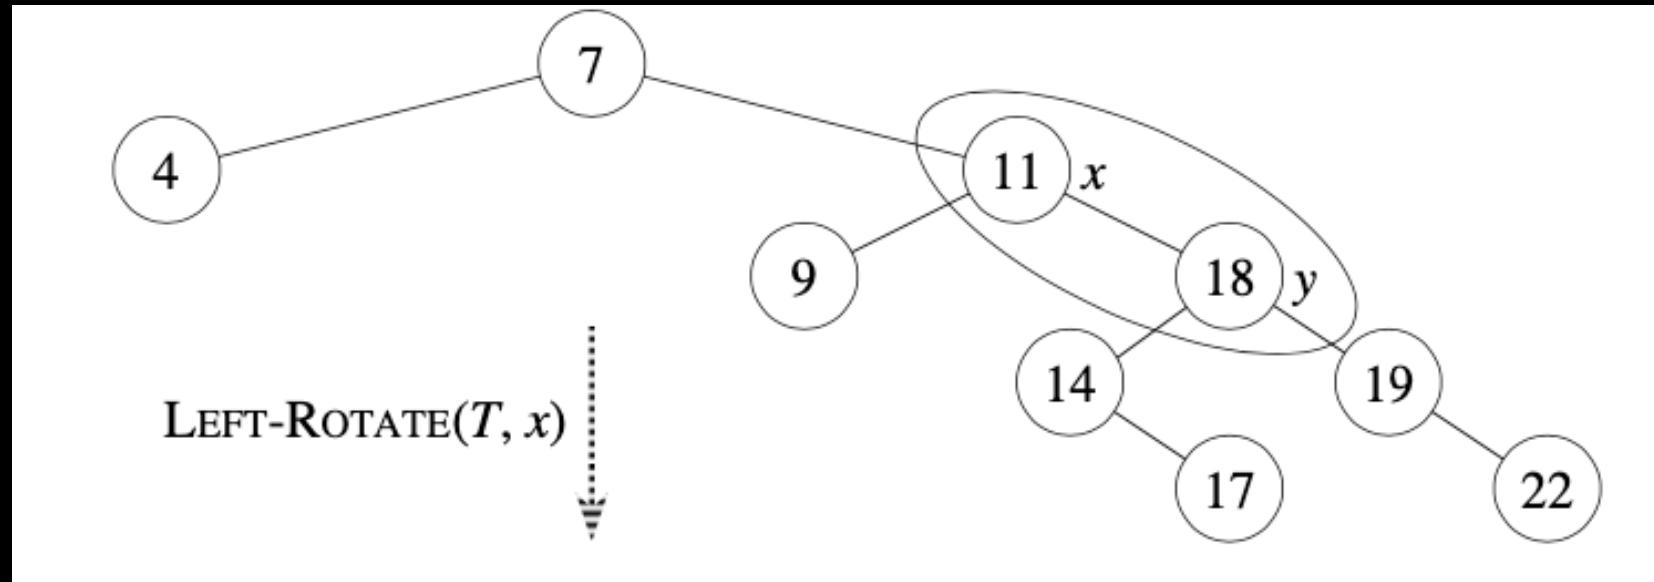
- (1) $y = x.\text{right}$ //set y
- (2) $x.\text{right} = y.\text{left}$ //turn left subtree of y into right of x
- (3) if ($y.\text{left} \neq T.\text{nil}$) then
- (4) $y.\text{left}.p = x$
- (5) $y.p = x.p$ //link parent of x to y
- (6) if ($x.p = T.\text{nil}$) then
- (7) $T.\text{root} = y$
- (8) else
- (9) if ($x = x.p.\text{left}$) then
- (10) $x.p.\text{left} = y$
- (11) else
- (12) $x.p.\text{right} = y$
- (13) $y.\text{left} = x$ //put x of left of y
- (14) $x.p = y$



7.3.1. Rotations

- The pseudocode for LEFT-ROTATE assumes that
 - $x.\text{right} \neq T.\text{nil}$
 - root's parent is $T.\text{nil}$.
- The running time is $T(n) = O(1)$.
- The RIGHT-ROTATE algorithm is symmetric:
 - exchange left and right everywhere.

7.3.1. Rotations





7.3.2. Insertion

Insertion Consideration:

- Suppose a node z is inserted into RBT.
- If $z.\text{color} = \text{Red}$?
 - Property 1:
 - Property 2:
 - Property 3:
 - Property 4:
 - Property 5:



7.3. Operations

Insertion Consideration:

- Suppose a node z is inserted into RBT.
- If $z.\text{color} = \text{Black}$?
 - Property 1:
 - Property 2:
 - Property 3:
 - Property 4:
 - Property 5:



7.3.2. Insertion

- Recall BST insertion:
 - Two pointers, x and y, were used to find the location.
 - A new node z will be inserted as BST insertion. Why?
 - z.color?
 - We need an additional algorithm to fix to maintain RBT color properties.



7.3.2. Insertion

Algorithm (TREE-INSERT(T, x))

- (1) $y = \text{NIL}$, $x = T.\text{root}$
- (2) while ($x \neq \text{NIL}$) do
- (3) $y = x$
- (4) if ($z.\text{key} < x.\text{key}$) then
- (5) $x = x.\text{left}$
- (6) else $x = x.\text{right}$
- (7) $z.p = y$
- (8) if ($y = \text{NIL}$) then
- (9) $T.\text{root} = z$
- (10) else if ($z.\text{key} < y.\text{key}$) then
- (11) $y.\text{left} = z$
- (12) else $y.\text{right} = z$



7.3.2. Insertion

- The color properties are violated when $z.\text{color} = z.p.\text{color} = \text{red}$.
- The new balancing considers the following cases depending on the color of z 's uncle, $y.\text{color}$.
 1. $z.p = z.p.p.\text{left}$ & $y.\text{color} = \text{red}$ (case 1)
 1. $z = z.p.\text{right}$ & $y.\text{color} = \text{black}$ (case 2)
 2. $z = z.p.\text{left}$ & $y.\text{color} = \text{black}$ (case 3)
 2. Symmetric cases when $z.p = z.p.p.\text{right}$.

7.3.2. Insertion



Case 1: $z.p = z.p.p.left$ and $y.color = \text{red}$

- Assume $z.p.p.color = \text{black}$.
- **Make $z.p.color = y.color = \text{black}$.**
 - Property 4:
 - Property 5:
- **Make $z.p.p.color = \text{red}$.**
- Let new z be $z.p.p$.

7.3.2. Insertion



Case 2: $z = z.p.right$ and $y.color = \text{black}$

- Let $z = z.p$ then
- **LEFT-ROTATE(T, z)**
 - Property 4:
 - Property 5:
- Transforms to case 3: $z = z.p.left$

7.3.2. Insertion



Case 3: $z = z.p.left$ and $y.color = \text{black}$

- **Make $z.p.color = \text{red}$ and $z.p.p.color = \text{black}$**
- **RIGHT-ROTATE($T, z.p.p$)**
 - Property 4:
 - Property 5:
- Transforms to case 3: $z = z.p.left$

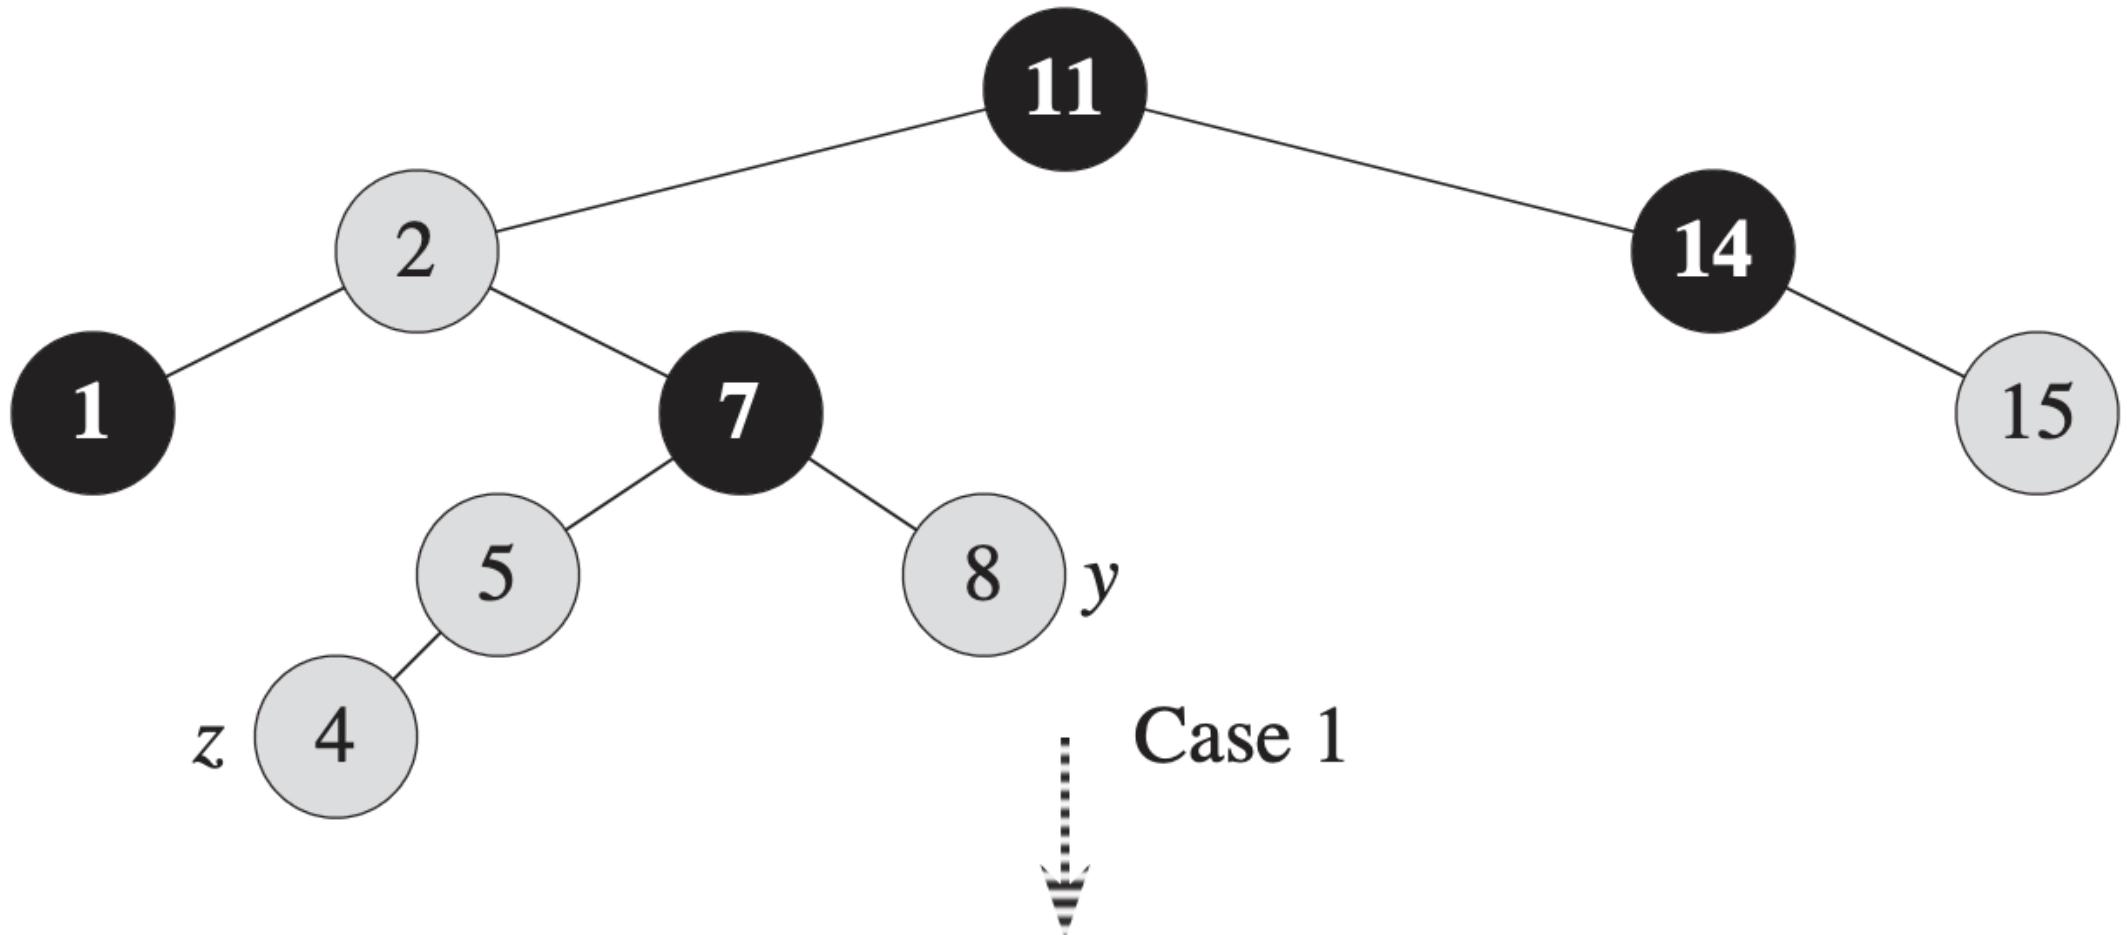


7.3.2. Insertion

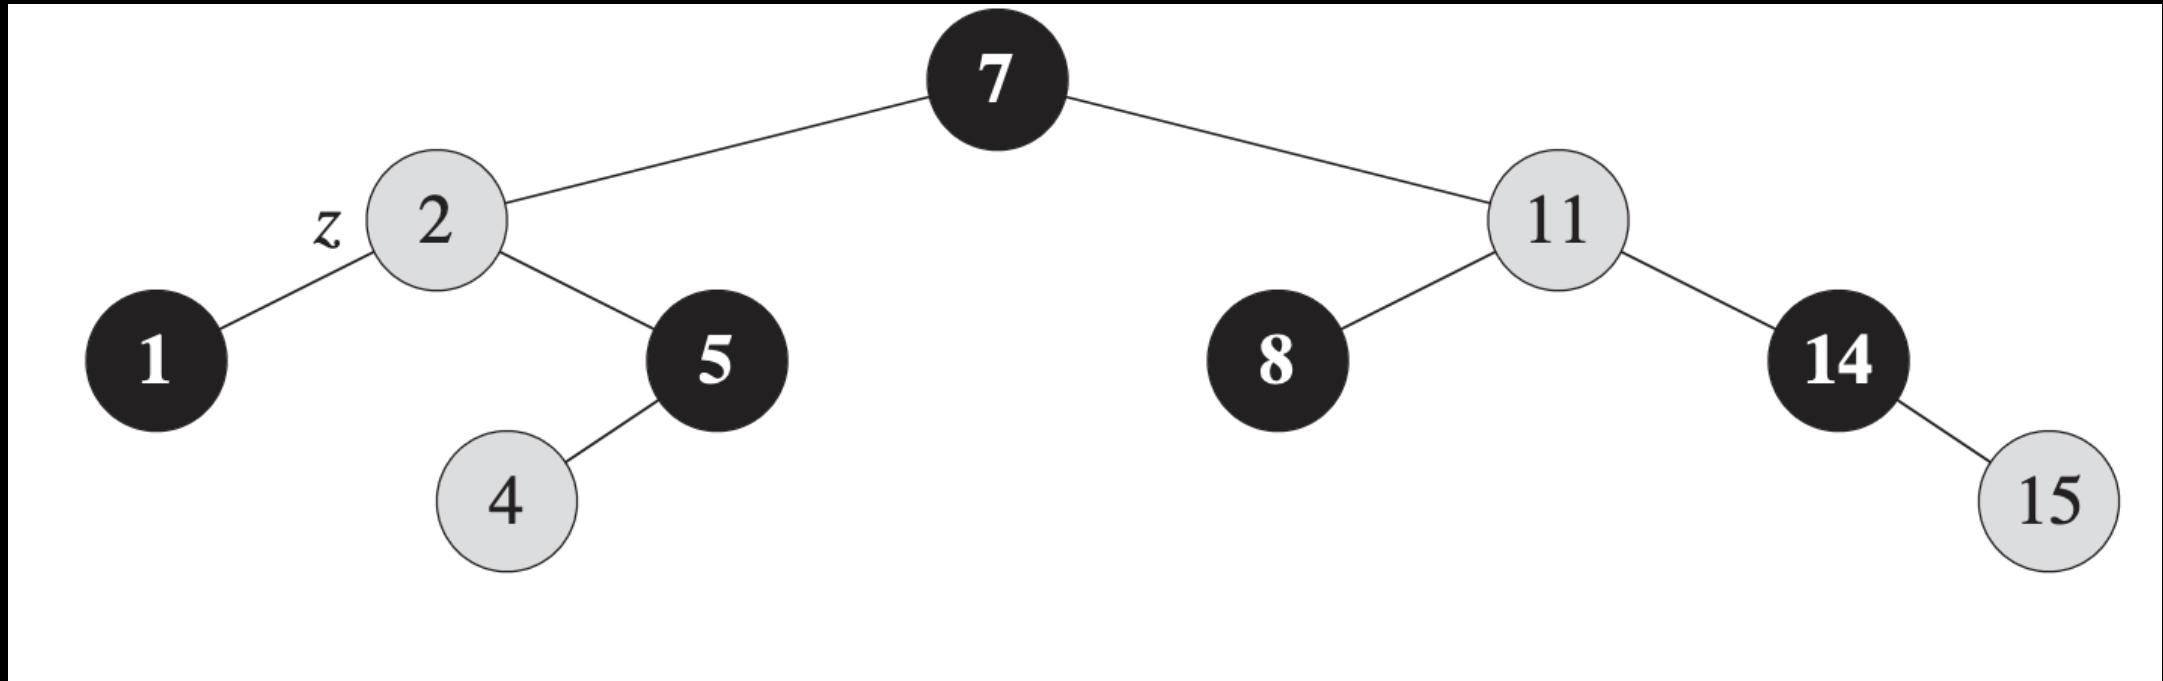
Algorithm (RB-INSERT-FIXUP(T, z))

```
(1)      while ( $z.p.color = \text{RED}$ )
(2)          if ( $z.p = z.p.p.left$ )
(3)               $y = z.p.p.right$ 
(4)              if ( $y.color = \text{RED}$ )
(5)                   $z.p.color = \text{BLACK}$ ,  $y.color = \text{BLACK}$  //case 1
(6)                   $z.p.p.color = \text{RED}$ 
(7)                   $z = z.p.p$ 
(8)              else
(9)                  if ( $z = z.p.right$ ) then
(10)                       $z = z.p$                                 //case 2
(11)                      LEFT-ROTATE( $T, z$ )
(12)                       $z.p.color = \text{BLACK}$ ,  $z.p.p.color = \text{RED}$  //case 3
(13)                      RIGHT-ROTATE( $T, z.p.p$ )
(14)              else
(15)                  (do same with right and left exchange)
(16)       $T.root.color = \text{BLACK}$ 
```

7.3.2. Insertion



7.3.2. Insertion





7.3.2. Insertion

Analysis

$O(\lg n)$ time to get through RB-INSERT up to the call of RB-INSERT-FIXUP.

Within RB-INSERT-FIXUP:

- Each iteration takes $O(1)$ time.
- Each iteration is either the last one or it moves z up 2 levels.
- $O(\lg n)$ levels $\Rightarrow O(\lg n)$ time.
- Also note that there are at most 2 rotations overall.

Thus, insertion into a red-black tree takes $O(\lg n)$ time.



7.3. Operations

Deletion Consideration:

- Suppose a node z is removed.
 - If $z.\text{color} = \text{red}$?
 - Property 1:
 - Property 2:
 - Property 3:
 - Property 4:
 - Property 5:



7.3. Operations

Deletion Consideration:

- Suppose a node z is removed.
 - If $z.\text{color} = \text{black}$?
 - Property 1:
 - Property 2:
 - Property 3:
 - Property 4:
 - Property 5:

7.3.3. Deletion



- Recall BST deletion operations:
 - z having a single child:
 - z with two children:
- In RBT deletion,
 - Similar to BST deletion, the transplant algorithm will be used.
 - Case operations from BST deletion will be carried over.

7.3.3. Deletion



- These are the differences from the BST deletion:
 1. A node y will be used as a removing node.
 2. $y.\text{color}$ must be identified at the beginning and saved.
 1. y 's original color will be saved as $y\text{-original-color}$.
 3. x will replace y (BST deletion replaces z with y)
 4. Call $\text{RB_DELETE_FIXUP}(T, x)$ if $y\text{-original-color} = \text{black}$.



7.3.3. Deletion

(1)	RB-TRANSPLANT(T, u, v) if u.p == T.nil T.root = v	RB-DELETE(T, z) y = z y-original-color = y.color
(2)	elseif u == u.p.left u.p.left = v	if z.left == T.nil x = z.right RB-TRANSPLANT(T, z, z.right)
(3)	else u.p.right = v	elseif z.right == T.nil x = z.left RB-TRANSPLANT(T, z, z.left)
(4)	v.p = u.p	else y = TREE-MINIMUM(z.right) y-original-color = y.color x = y.right if y.p == z x.p = y
(5)		else RB-TRANSPLANT(T, y, y.right) //y in subtree of z subtree
(6)		y.right = z.right //and not root of subtree
(7)		y.right.p = y
(8)		RB-TRANSPLANT(T, z, y) // replace z by y
(9)		y.left = z.left
(10)		y.left.p = y
(11)		y.color = z.color
(12)		if y-original-color == BLACK RB-DELETE-FIXUP(T, x)
(13)		
(14)		
(15)		
(16)		
(17)		
(18)		
(19)		
(20)		
(21)		
(22)		

// z has no left child

// just left child

//two children, y is successor of z

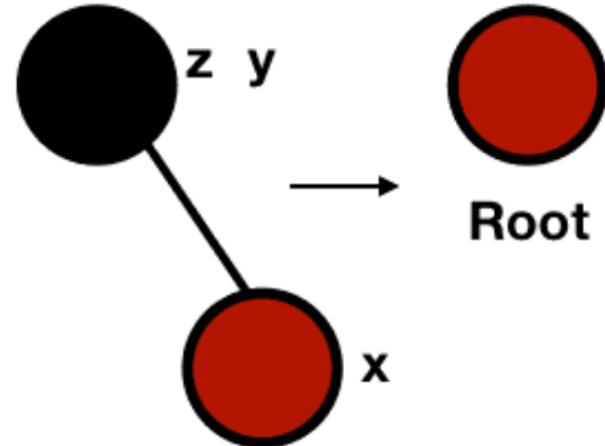
- Deleting a node from a red-black tree is a bit more complicated than inserting a node.
- The procedure for deleting a node from a red-black tree is based on the TREE-DELETE procedure.



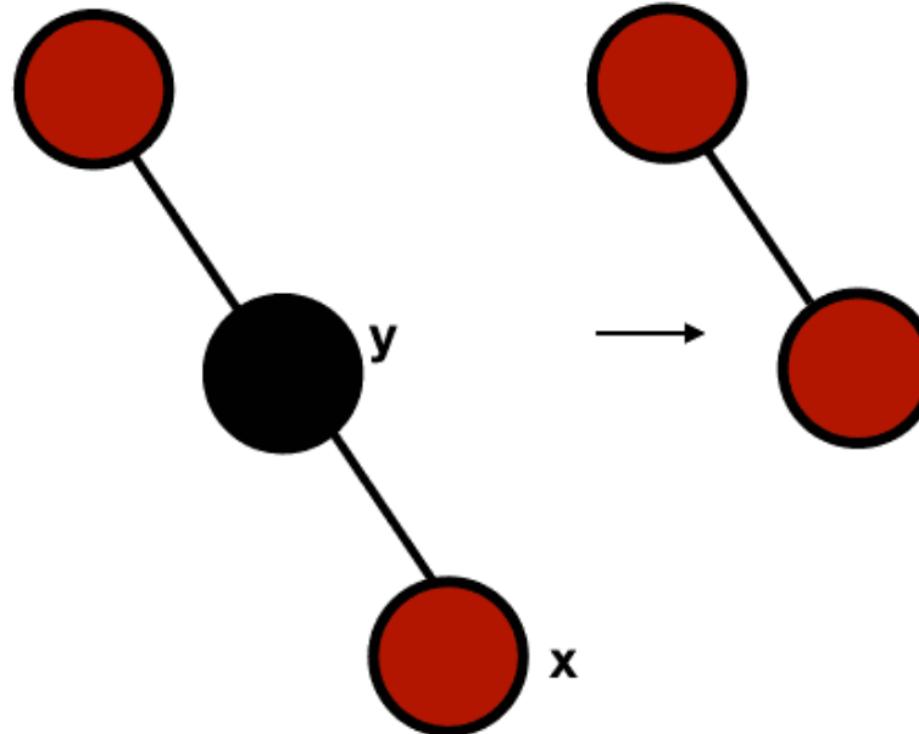
7.3.3. Deletion

- When the RBT 5th property is violated,
 - Give x an “extra black”.
 - x is either doubly black (if x.color=BLACK) or red & black (if x.color=RED).
 - Property 5 satisfied, but property 1 violated.
 - The attribute x.color is still either RED or BLACK. No new values for the color attribute.
 - Extra blackness on a node is by virtue of x pointing to the node.

7.3.3. Deletion



Violation of prop. 2



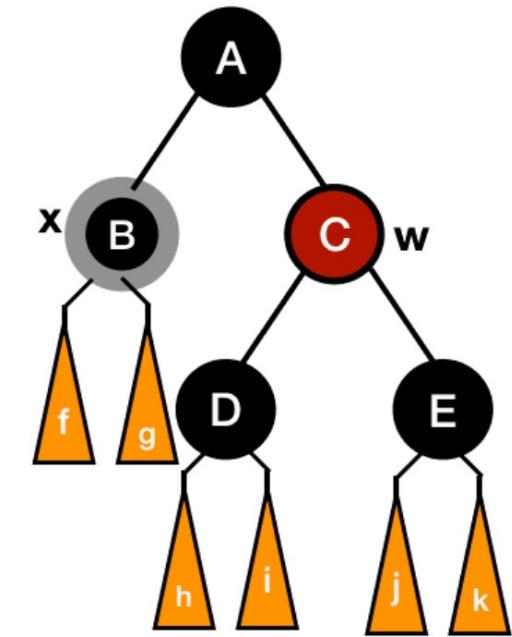
Violation of prop. 4

**Violation of prop. 5,
black height affected**

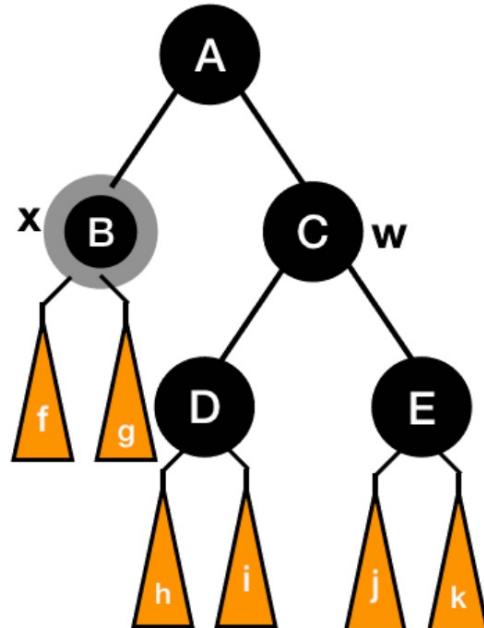
7.3.3. Deletion



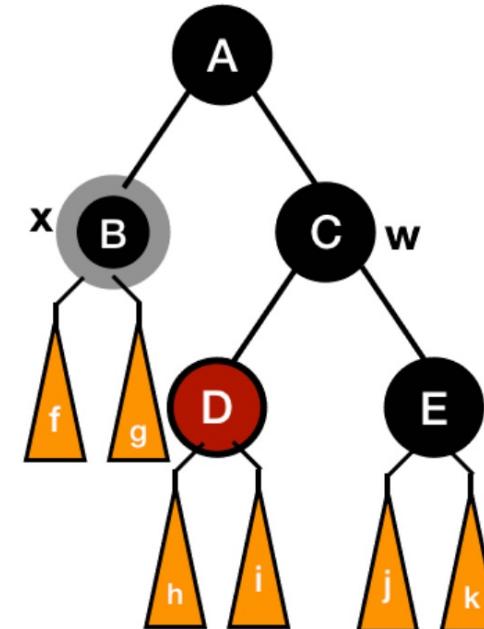
- There are four cases to consider after the deletion when $x = x.p.left$



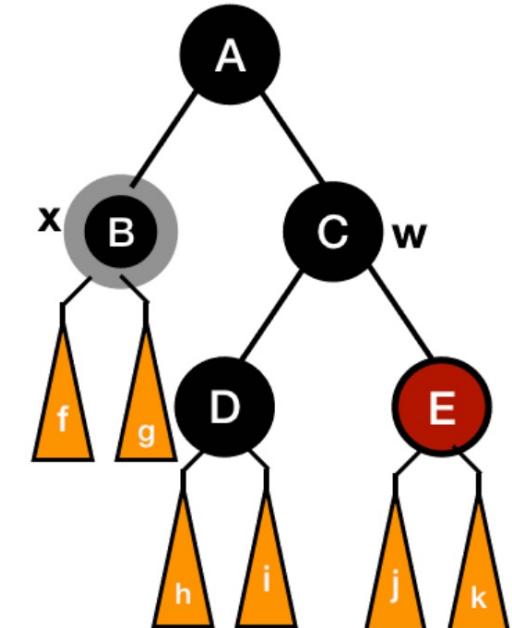
Case 1



Case 2

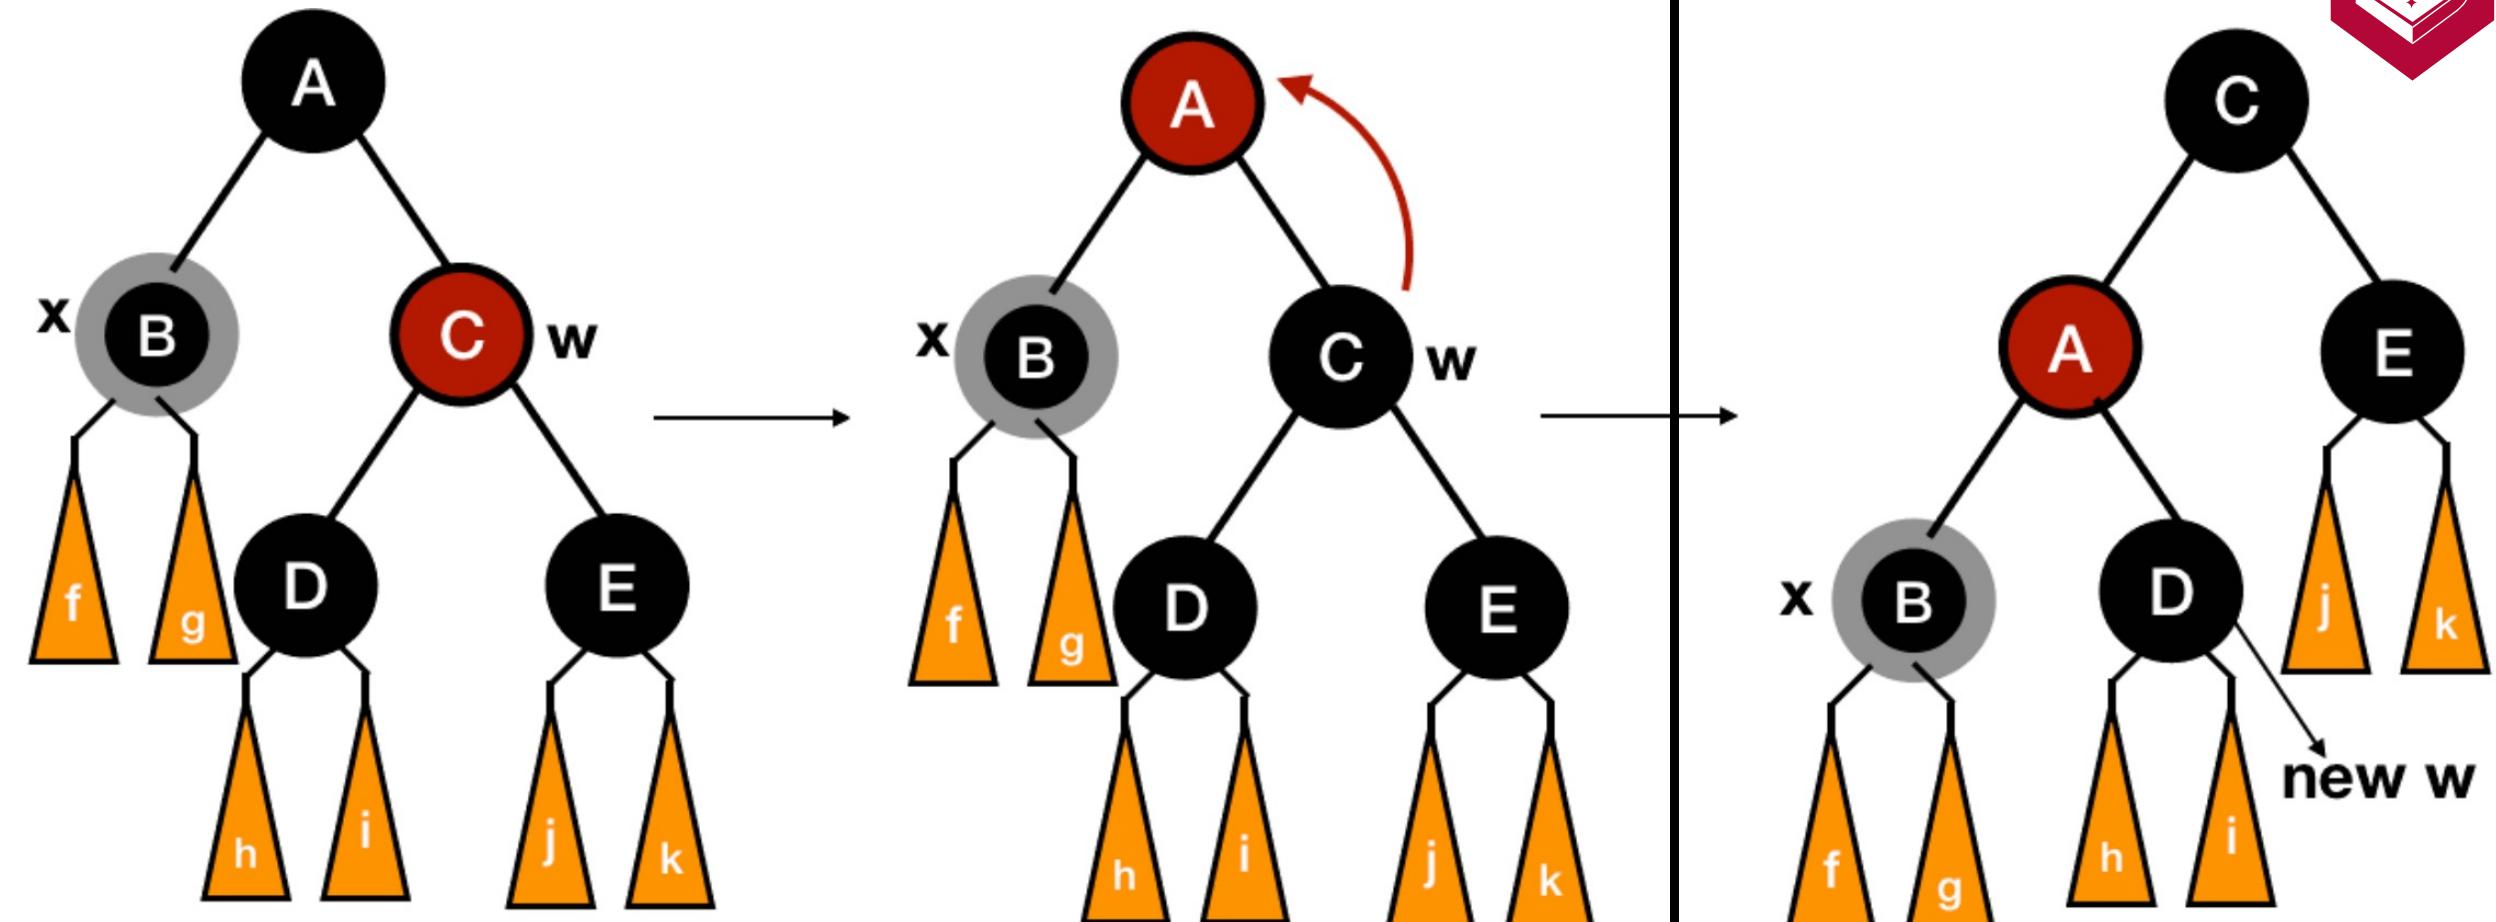


Case 3



Case 4

7.3.3. Deletion



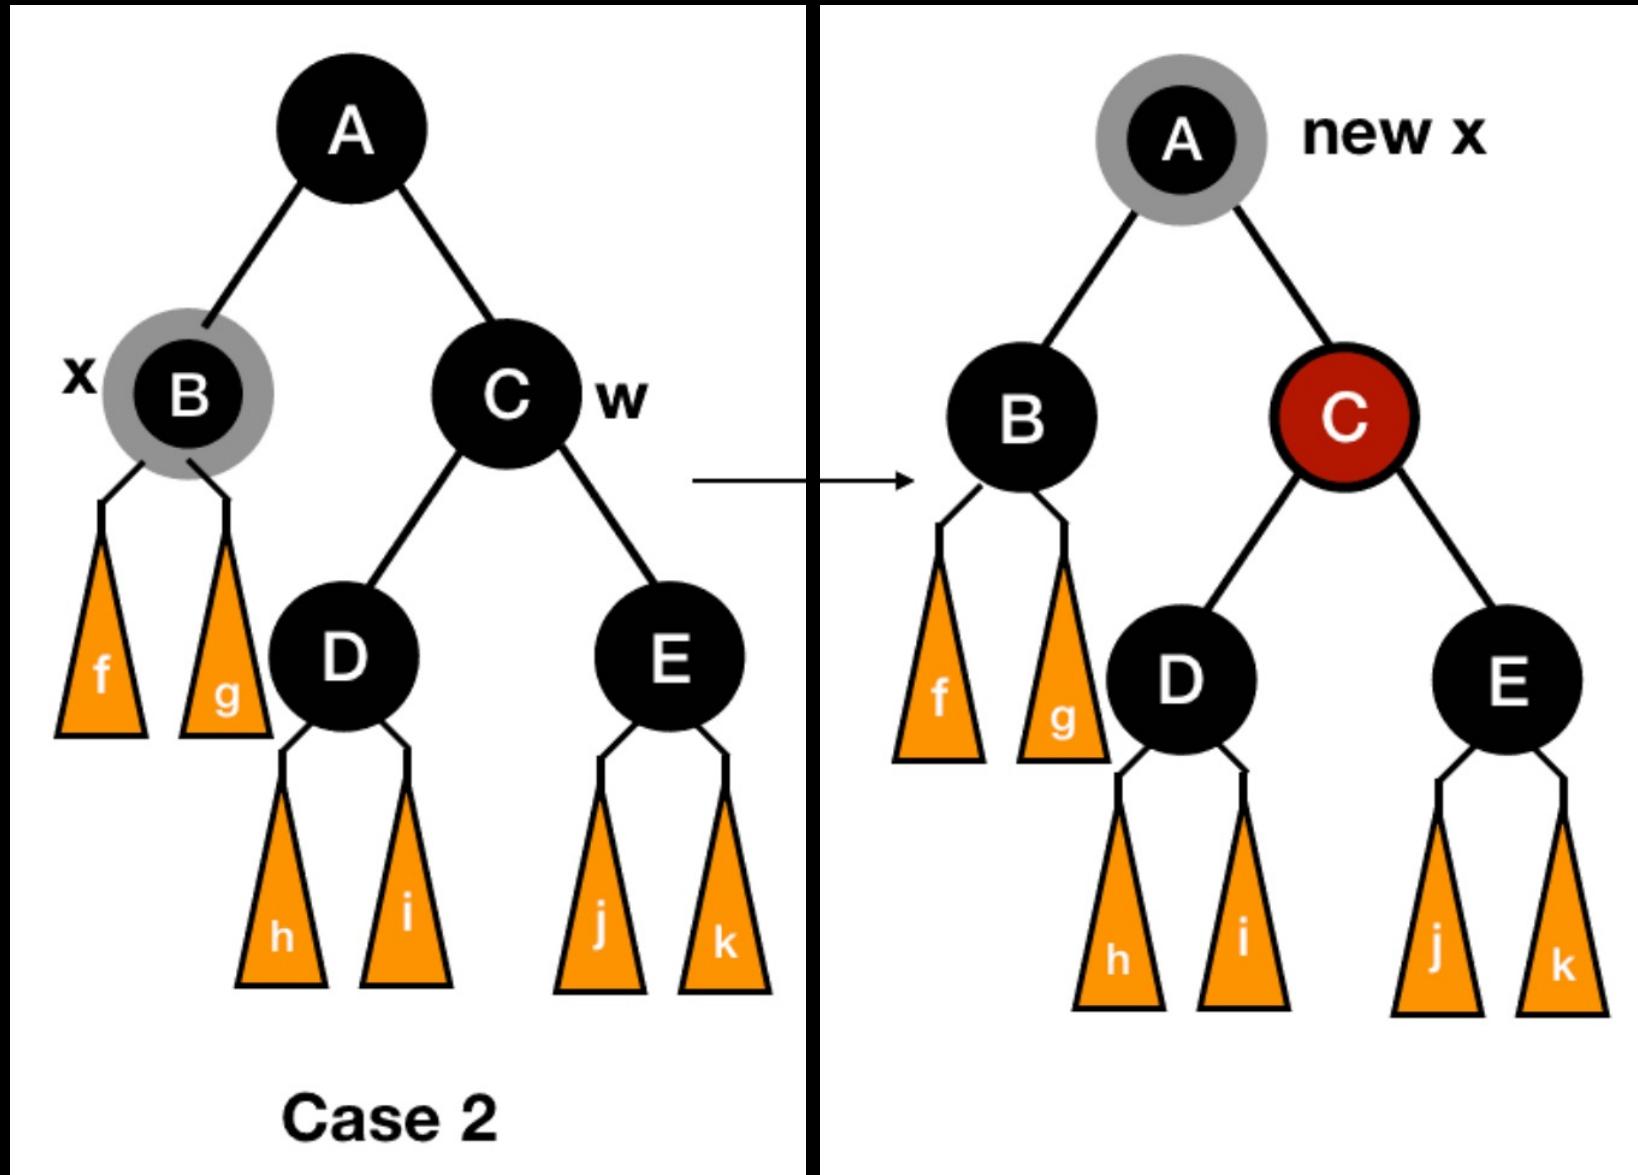
Case 1



7.3.3. Deletion

- $w.left.color = w.right.color = \text{black}$ and $w.color = \text{red}$.
- Make $w.color = \text{black}$ and $x.p.color = \text{red}$.
- Then $\text{LEFT_ROTATE}(T, x.p)$.
- w 's child is a new sibling of x
- Make a new $w = x.p.right$
- Go immediately to cases 2, 3, or 4.

7.3.3. Deletion

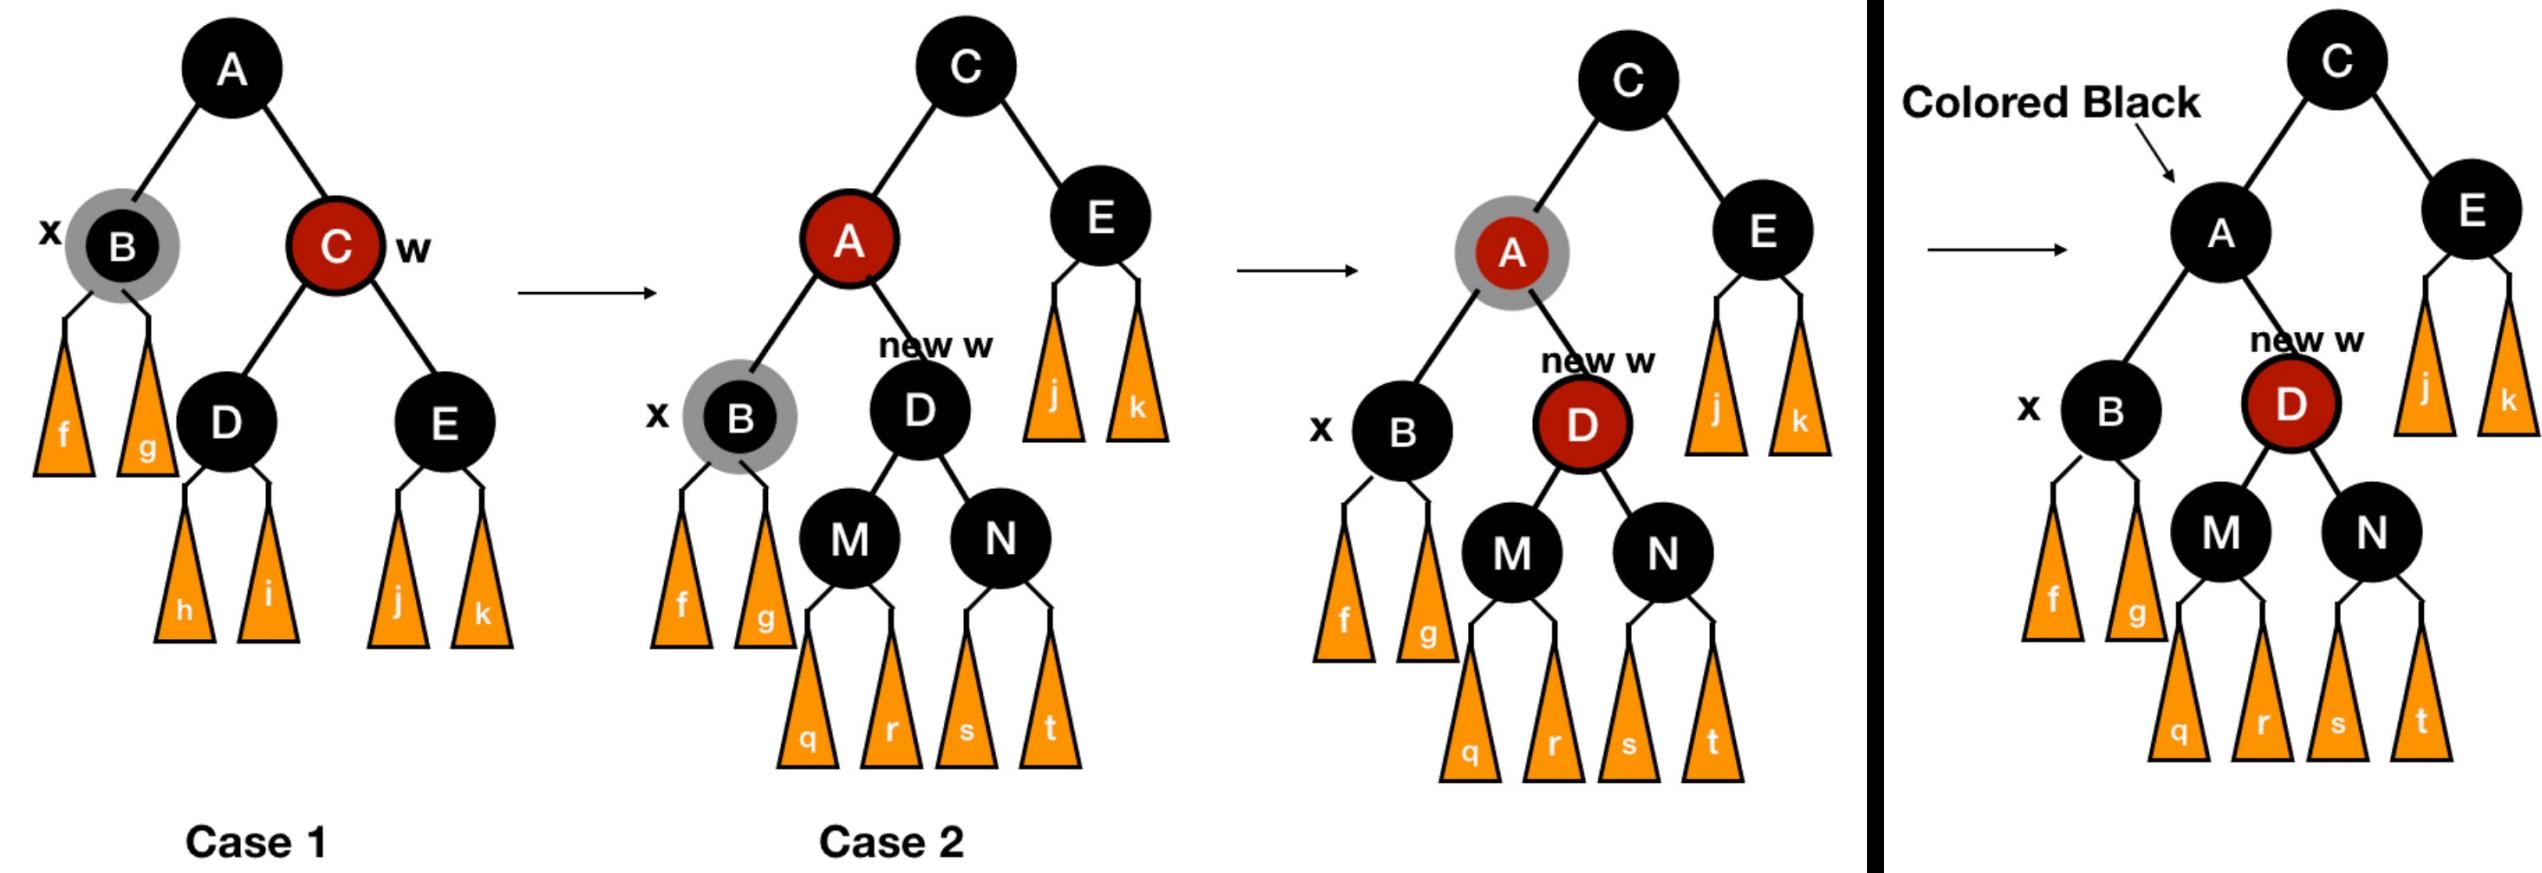


7.3.3. Deletion

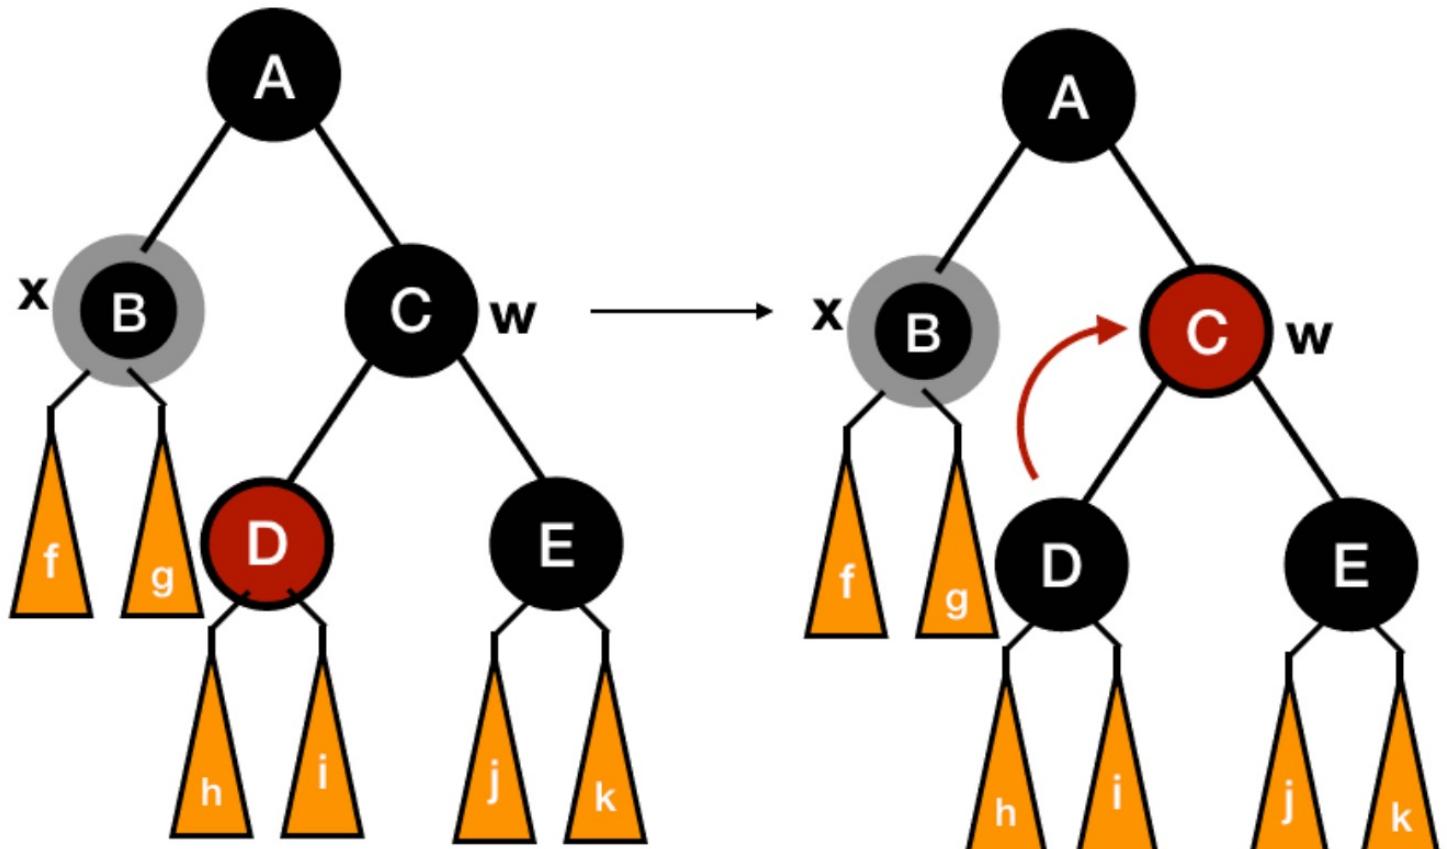


- Case 2: if $w.\text{left}.\text{color} = w.\text{right}.\text{color} = w.\text{color} = \text{black}$
 - Make $w.\text{color} = \text{red}$
 - Move up to have $x = x.p$

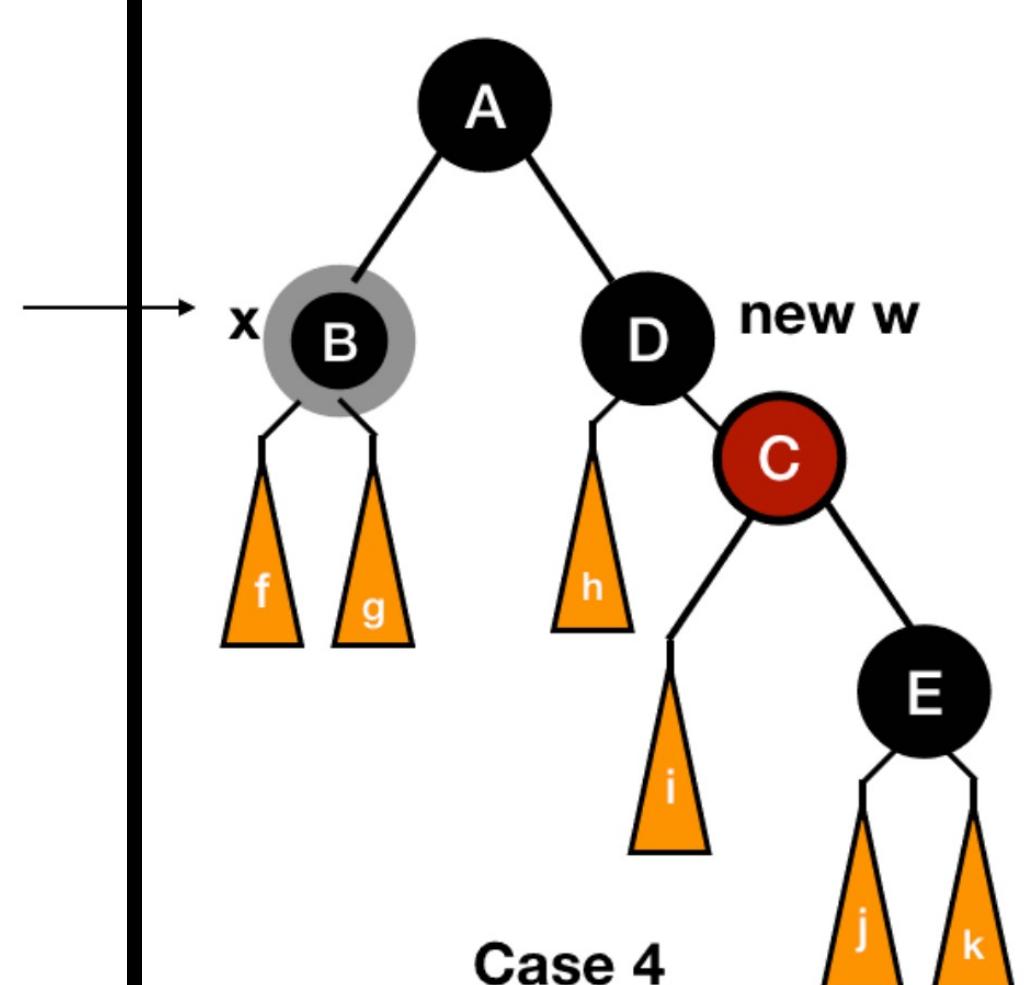
7.3.3. Deletion



7.3.3. Deletion



Case 3



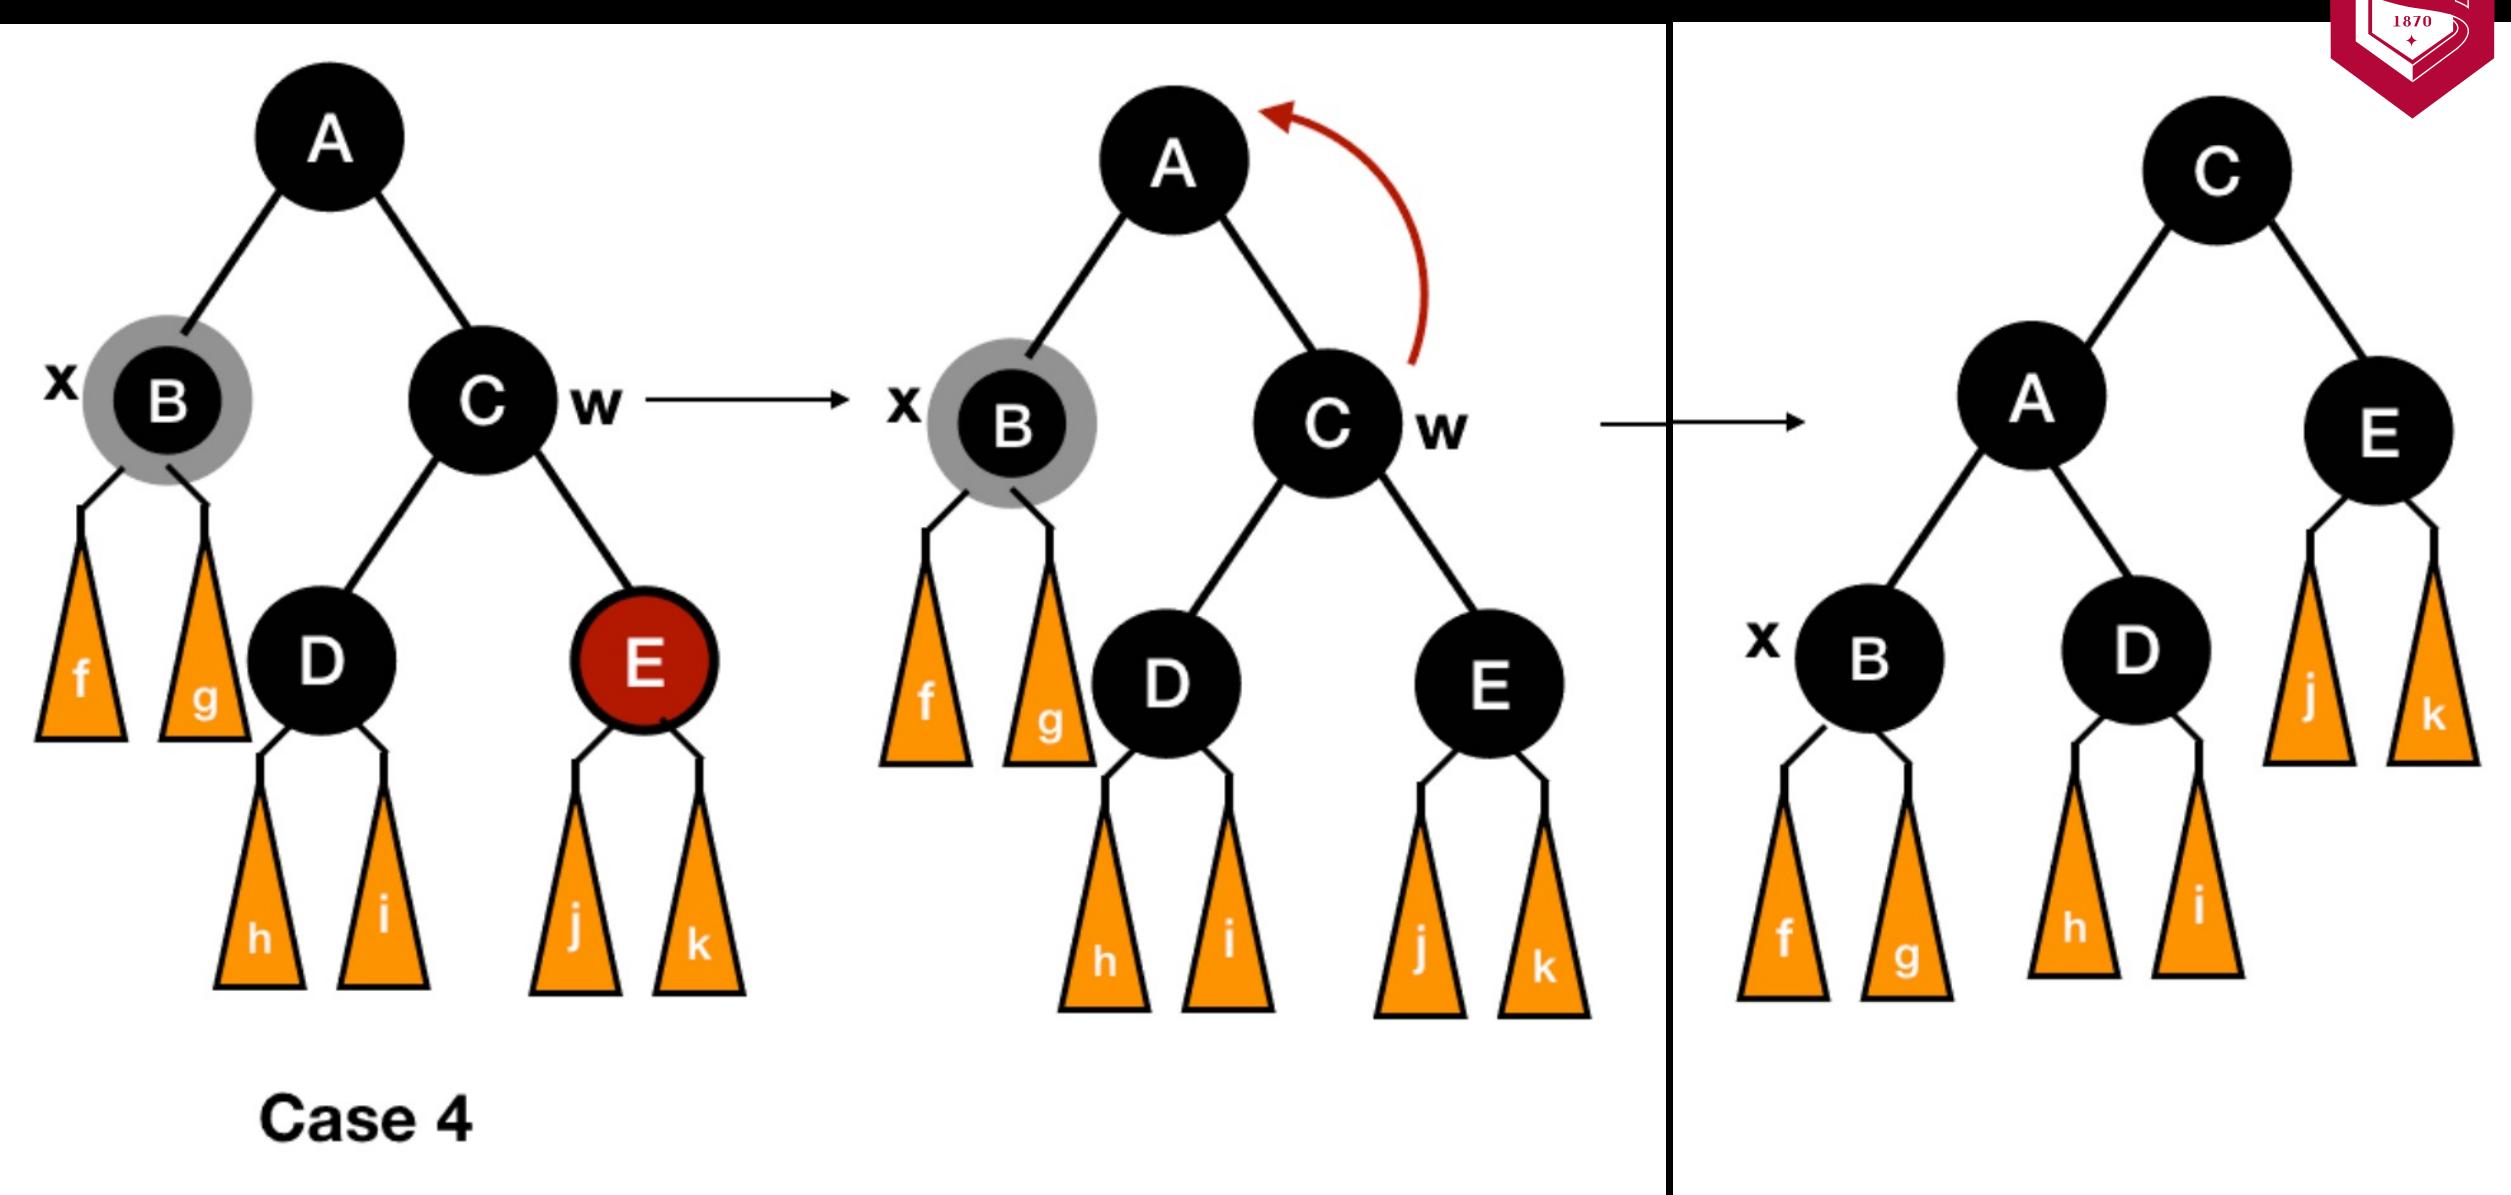
Case 4

7.3.3. Deletion



- Case 3: $w.\text{color} = w.\text{right}.\text{color} = \text{black}$ and $w.\text{left}.\text{color} = \text{red}$
- Make $w.\text{color} = \text{red}$ and $w.\text{left}.\text{color} = \text{black}$
- Then, **RIGHT_ROTATE(T, w)**
- x 's new sibling $w.\text{color} = \text{black}$ and $w.\text{right}.\text{color} = \text{red}$

7.3.3. Deletion





7.3.3. Deletion

- Case 4: $w.\text{color} = w.\text{left}.\text{color} = \text{black}$ and $w.\text{right}.\text{color} = \text{red}$
- Make $w.\text{color} = x.p.\text{color}$, $x.p.\text{color} = \text{black}$,
 $w.\text{right}.\text{color} = \text{black}$
- Then, $\text{LEFT_ROTATE}(T, x.p)$.
- Remove extra black on x .
- Set $x = T.\text{root}$
- All done.



7.3.3. Deletion

```
RB-DELETE-FIXUP(T,x)
(1) while (x ≠ T.root and x.color == BLACK) do
(2)     if (x == x.p.left) then
(3)         w = x.p.right
(4)         if (w.color == RED) then
(5)             w.color == BLACK          //case 1
(6)             x.p.color = RED
(7)             LEFT-ROTATE(T, x.p)
(8)             w = x.p.right
(9)         if (w.left.color == BLACK and w.right.color == BLACK) then
(10)             w.color = RED           //case 2
(11)             x = x.p
(12)         else if (w.right.color == BLACK) then
(13)             w.left.color = BLACK    //case 3
(14)             w.color = RED
(15)             RIGHT-ROTATE(T,w)
(16)             w = x.p.right
(17)             w.color = x.p.color   //case 4
(18)             x.p.color = BLACK
(19)             w.right.color = BLACK
(20)             LEFT-ROTATE(T, x.p)
(21)             x = T.root
(22)         else (same as then clause with "right" and "left" exchanged)
(23)         x.color = BLACK
```



7.3.3. Deletion

Idea: Move the extra black up the tree until

- x points to a red & black node \Rightarrow turn it into a black node,
- x points to the root \Rightarrow remove the extra black, or
- we can do certain rotations and re-colorings and finish.

Within the **while** loop:

- x always points to a non-root doubly black node.
- w is x 's sibling.
- w cannot be $T.nil$, since that would violate property 5 at $x.p$.

There are 8 cases, 4 of which are symmetric to the other 4. As with insertion, the cases are not mutually exclusive. We'll look at cases in which x is a left child.

7.3.3. Deletion



Analysis

$O(\lg n)$ time to get through RB-DELETE up to the call of RB-DELETE-FIXUP.

Within RB-DELETE-FIXUP:

1. Case 2 is the only case in which more iterations occur.
 - x moves up 1 level.
 - Hence, $O(\lg n)$ iterations.
2. Each case 1, 3, and 4 have 1 rotation $\Rightarrow \leq 3$ rotations.
3. Hence, $O(\lg n)$ time.